


LunarLander-v2: sensitivity to learning rate (Double DQN)



Learning rate was selected because it controls the scale of Adam updates and allows us to compare learning speed and stability of bootstrapped Q-learning. We compare the default $1e-3$ with $1e-4$, $3e-4$, and $3e-3$, keeping other training settings unchanged. Each run uses `--seed 1` and 500,000 steps; each point averages 10 greedy evaluation episodes. Curves are unsmoothed. The shaded region marks the last 10 logged evaluations (400,000-490,000 steps). This is a single-run comparison per setting, not an average over training seeds.